

General Sum Markov Games for Strategic Detection of Advanced Persistent Threats using Moving Target Defense in Cloud Networks

Sailik Sengupta, Ankur Chowdhary, Dijiang Huang, and
Subbarao Kambhampati

Arizona State University, Tempe, AZ, USA
{sailiks,achaud16,dijiang,rao}@asu.edu

Abstract. The processing and storage of critical data in large-scale cloud networks necessitate the need for scalable security solutions. It has been shown that deploying all possible detection measures incur a cost on performance by using up valuable computing and networking resources, thereby resulting in Service Level Agreement (SLA) violations promised to the cloud-service users. Thus, there has been a recent interest in developing Moving Target Defense (MTD) mechanisms that helps to optimize the joint objective of maximizing security while ensuring that the impact on performance is minimized. Often, these techniques model the real challenge of multi-stage attacks by stealthy adversaries as a single-step attack detection game using graph connectivity measures as a heuristic to measure performance, thereby (1) losing out on valuable information that is inherently present in multi-stage models designed for large cloud networks, and (2) come up with strategies that have asymmetric impacts on performance, thereby heavily affecting the Quality of Service (QoS) for some cloud users. In this work, we use the attack graph of a cloud network to formulate a general-sum Markov Game and use the Common Vulnerability Scoring System (CVSS) to come up with meaningful utility values in each state of the game. We then show that, for a threat model where an adversary has the knowledge of a defender’s strategy in a state, the use of Stackelberg equilibrium can provide an optimal strategy for security resource placement. In cases where this assumption threat model turns out to be too strong, we show that the Stackelberg equilibrium turns out to be a Nash equilibrium of the general-sum Markov Game in most cases. We compare the gains obtained by using our method(s) to other baseline techniques used in Markov games and in cloud network security. Finally, we highlight how the method was used in a real-world small-scale cloud system.

1 Introduction

A cloud service provider provides processing and storage hardware along with networking resources to customers for profit. Although a cloud provider might want to use state-of-the-art security protocols, vulnerabilities in software desired (or used) by customers can put sensitive information stored in or communicated over the cloud network at risk. Distributed elements such as firewalls, Intrusion Detection Systems (IDS), log-monitoring systems etc. have been the backbone

for detecting malicious traffic in (or stopping them from entering) such systems. Unfortunately, the scale of modern-day cloud systems makes the placement of all possible detecting and monitoring mechanisms an expensive solution [11, 31, 26]; using up the computing and network resources that could have been better utilized by giving them to customers which in turn would be better for business. Thus, the question of how one should place a limited number of detection mechanisms to limit the impact on performance while ensuring that the security of the system is not drastically weakened becomes a significant one.

There has been an effort to answer this question in previous research works [31, 26]. Researchers have pointed out that static placement of detection systems is doomed to be insecure because an attacker, with reconnaissance on their side (which exists for any such cyber-system by default), will eventually learn this static placement strategy and hence, avoid it. Thus, a dynamic placement of these detection mechanisms, generally known as Moving Target Defense for continuously shifting the detection surface, has become a default. In this method, the set of attacks for which monitoring systems are placed changes in some randomized fashion after every fixed time step. In these works, the cloud system is treated in a way similar to that of physical security systems where the primary challenge is to allocate a limited set of security resources to an asset/schedule that needs to be protected [22, 30].

In the case of real-world cloud-systems, these aforementioned solutions lead to three problems. First, only single-step attacks are considered in the game-theoretic modeling which lead to sub-optimal strategies because such models fail to account for multi-stage attack behavior. For example, strategies generated by prior work may prioritize detecting a high-impact attack on a web-server more than a low-impact attack on a path that leads an attack on a storage server, which when exploited may have major consequences. Second, the threat model assumes that an attacker can launch an attack from any node in the cloud network. This is too strong an assumption, leading to sub-optimal placement strategies in real-world settings. Third, existing methods can come up with placement strategies that allocate multiple detection systems on a sub-net while another sub-net is not monitored. This results in a steep degradation of performance for some customers. To address these challenges, we need a suitable method for modeling multi-stage attacks. Capturing all possible attack paths can lead to an action-set explosion for previously proposed normal-form games. Thus, we use *Markov Games* to model such interactions.

Specifically, we try to address these problems by modeling the cloud system as a General-Sum Markov Game. We use particular nodes of our system's attack graph to represent the states of our game while the attacker's actions are modeled after real-world attacks based on the Common Vulnerabilities and Exploits (CVEs) described in the National Vulnerability Database (NVD) [1]. The defender's actions correspond to the placement of detection systems that can detect the attacks. We design the utility values for each player in this game leveraging (1) the Common Vulnerability Scoring Systems, which provide metrics for each attack, and (2) cloud designer's quantification of how the placement of a detection system impacts the performance. These help us come up with defense strategies that take into account the long-term impacts of a multi-stage attack while restricting the defender to pick a limited number of monitoring actions

in each part of the cloud. The latter constraints ensure that the performance impact on a customer, due to the placement of detection measures, is limited.

The popular notion of using min-max equilibrium for Markov Games [17] is an optimal strategy for players¹ only in zero-sum games and becomes sub-optimal for a general-sum game. Thus, we assume that an attacker is aware of the defender’s placement strategy at each state of our Markov game and consider the Stackelberg equilibrium of this game. In scenarios where this assumption becomes unrealistic, we show that the Stackelberg Equilibrium is, depending on the problem structure, a subset of Nash Equilibria and thereby, results in optimal strategies. The key contributions of this research work are as follows:

- We model the multi-stage attack scenarios, which are typically employed in launching Advanced Persistent Threats (APTs) campaigns against high-value targets, as a general-sum Markov Game. The cost-benefit analysis based on the two-player game, provides strategies for placing detection agents in a cloud network.
- We leverage the attack-graph modeling of cloud networks, the Common Vulnerabilities and Exposures (CVEs) present in the National Vulnerability Database and the Common Vulnerability Scoring Service (CVSS) to design the states, the actions and utility values of our game. In addition, we consider prior work in cloud-systems to (1) model the uncertainty of attack success and (2) leverage heuristic measures that model the performance impact of placing detection mechanisms on the cloud.
- Our framework considers a threat model where the attacker can infer the defender’s placement in a state of the Markov Game. Therefore, we design a dynamic programming solution to find the Stackelberg equilibrium of the Markov Game. In cases where this model is not an accurate representation of reality, we show that the Stackelberg equilibrium is a subset of Nash equilibrium when a set of properties hold (similar to prior work in extensive form games [15]). In order to showcase the effectiveness of our approach we analyze a synthetic and a real-world cloud system.

2 Background

In this section, we first introduce the reader to the notion of real-world vulnerabilities and exploits present in a cloud system that we will use throughout our paper. Second, we describe the threat model for our cloud scenario. Lastly, we describe the notion of Attack Graphs followed by a brief description of Markov games and some well-known algorithms to find the optimal policy or strategy of each player. In doing all this, we use an example attack scenario for cloud networks shown in Figure 1.

2.1 Vulnerabilities and Exploits

Software security is defined in terms of three characteristics - Confidentiality, Integration, and Availability [19]. Thus, in a broad sense, a vulnerability (that

¹ A preliminary version of this work was accepted to AAAI-19 workshop on Artificial Intelligence for Cyber Security and deals only with the zero-sum case.

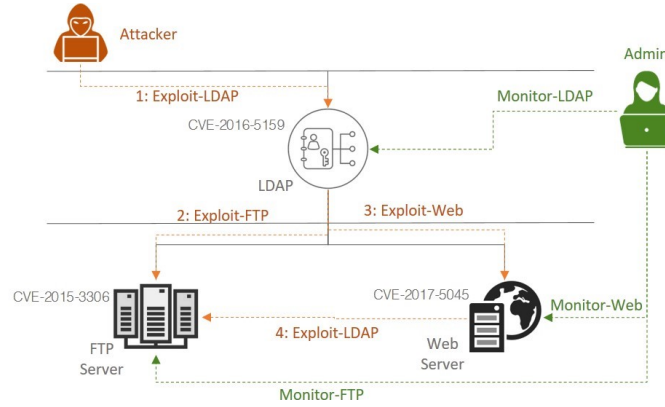


Fig. 1. An example cloud system highlighting its network structure, the attacker and defender (admin) agents and the possible attacks and monitoring mechanisms.

can be attacked or exploited) for a cloud system can be defined as a security flaw in a software service hosted over a given port, that when exploited by a malicious attacker, can cause loss of Confidentiality, Availability or Integrity (CIA) of that virtual machine (VM). The National Vulnerability Database (NVD) is a public directory of known vulnerabilities and exploits. It assigns an attack a unique identifier (CVE-id), describes the technology they affect and the attack behavior. Thus, to model the known attacks against our system, we use the Common Vulnerabilities and Exposures (CVEs) listed in this database.

In the cloud-scenario described in Figure 1, we have three VMs— an LDAP server, an FTP server, and a Web server. Each of these servers have a (set of) vulnerability present in it. On the LDAP server, an attacker can use a local privilege escalation to gain root privilege on it. The other two vulnerabilities— A cross-side scripting (XSS) attack on the Web server and the remove code execution on the FTP server— can only be executed with root access to the LDAP server. We can now describe the threat model for our scenario.

2.2 Threat Model

In the example cloud scenario, the attacker starts with user-level access to an LDAP server. Thus, this is the initial state for an attacker located outside the network. The terminal state is to compromise the FTP server (which, as we will see later, leads to an all absorbing state in our Markov Game). The attacker can perform actions such as 1: **exploit-LDAP**, **exploit-Web** or **exploit-FTP**. Note that in the scenario above the attacker has two possible paths to reach the goal node, i.e. **priv(attacker, (FTP: root))** which are:

- Path 1: **exploit-LDAP** → **exploit-FTP**
- Path 2: **exploit-LDAP** → **exploit-Web** → **exploit-FTP**

On the other hand, the (network) Admin, who is the defender in our case, can choose to monitor (1) read-write requests made by services running on a

VM using host-based Intrusion Detection Systems (IDS) like `auditd`, or (2) network traffic along both the paths using the network-based monitoring agents like `Snort`. We will denote these IDS systems using the terminology `monitor-LDAP`, `monitor-FTP`, etc. We assume that the Admin has a limited budget, i.e., cannot place all possible IDS system on the cloud network, and thus, must try to perform monitoring in an optimized fashion. On the other hand, the attacker will try to perform attacks along with some path that minimizes their probability of getting detected (and thus, maximizing their impact). Further, we assume an attacker has knowledge of the defender's placement strategy because of a reconnaissance phase it inherently has. Note that this makes pure strategies for placement of detection systems useless. For example, if the defender is only monitoring traffic only between LDAP and FTP server to detect an attack on the FTP server from the LDAP server, the attacker should ideally choose *Path 2* to avoid detection. Thus, to come up with a good dynamic placement strategy, we have to model the various multi-attack paths and the attacker's strategy. We first discuss the formalism of attack graphs that are a popular way to model the various attacks (and attack paths) in a cloud scenario [11, 4] and then move to a brief overview of two-player Markov Games that will help us reason about the attacker's behavior.

2.3 Attack Graph Formalism

Attack graphs (AG) are a representation tool to that help to model the security setting of a complex network system like the cloud. AGs have been shown to be useful in detecting multi-stage or multi-hop attack behavior which may not be obvious to a network administrator by (local) vulnerability scan analysis of the individual components (VMs or servers) part of the network.

Attack Graph is a graph $G = \{N, \mathcal{E}\}$ that consists of a set of nodes (N) and a set of edges (\mathcal{E}) where,

- As shown in the Figure 2, nodes can be of four types– the nodes N_C represent vulnerabilities (shown as rectangles), e.g. `vulExists (LDAP, Local Priv. Escalation)`, N_D represents the attacker's state (shown as diamonds) e.g., `priv(attacker, (LDAP :user))`, rule nodes N_R represent a particular exploit action (shown as ellipses) and finally, root or goal nodes N_G that represent the goal of an attack scenario (shown using two concentric diamonds), e.g., `priv(attacker, (FTP: root))`.
- $E = E_{pre} \times E_{post}$ denotes a set of directed edges in the graph. An edge $e \in E_{pre}$ goes from a node in N_D or N_C to a rule node N_R and denotes that an attack or rule can only be triggered if all the conditions of the edges going into $n \in N_R$ is satisfied (AND-nodes). An edge $e \in E_{post}$ goes from a node N_R to an attacker's state $n \in N_D$ and indicates that when a rule is executed, the attacker's privilege changes.

Note how the two attacks paths mentioned in the threat model section become evident by a simple look at the attack graph. The conditional and cumulative probability values for AND (conjunct) and OR (disjunct) nodes are calculated using probability estimates as described by Chung *et. al.* [5].

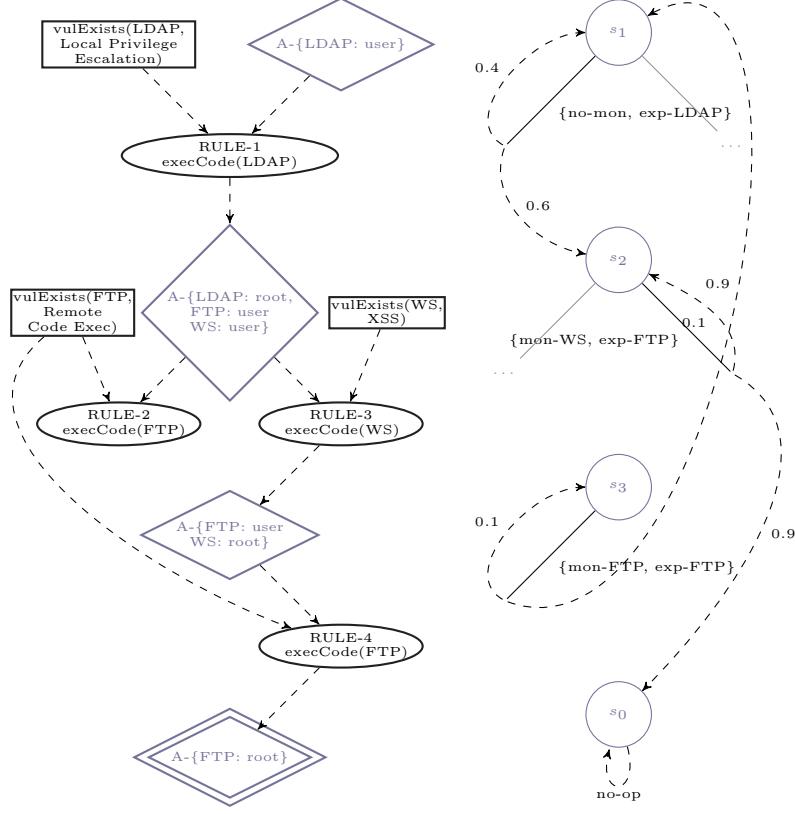


Fig. 2. The left figure shows the attack graph of the synthetic cloud scenario shown in 1. The right figure shows the formulated Markov Game.

2.4 Two-Player Markov Games

We use this opportunity of defining a Markov Game to also introduce the reader to the notations used in our modeling. The two players will be called the Defender D (who is the Admin of the cloud system) and the Attacker A (who is the adversary trying to exploit a vulnerability in the Cloud System). With that, we can now define a Markov Game as follows.

Markov Game for two players D and A can be defined by the tuple $(S, M, E, \tau, U^D, U^A, \gamma^D, \gamma^A)$ [29] where,

- $S = \{s_1, s_2, s_3, \dots, s_k\}$ are finite states of the game,
- $M = \{m_1, m_2, \dots, m_n\}$ represents the finite set of (monitoring) actions available to D ,
- $E = \{e_1, e_2, \dots, e_{n'}\}$ represents the finite set of (exploit) actions available to A ,

VM	Vulnerability	CVE	CIA Impact	Attack Complexity
LDAP	Local Priv Esc	CVE-2016-5195	5.0	MEDIUM
Web Server (WS)	Cross Site Scripting	CVE-2017-5095	7.0	EASY
FTP	Remote Code Exec.	CVE-2015-3306	10.0	MEDIUM

Table 1. Vulnerability Information for the Cloud Network

- $\tau(s, m_i, e_j, s')$ represents the probability of reaching a state $s' \in S$ from the state $s \in S$ if D choose to deploy the monitoring m_i and A choose to use the exploit e_j ,
- $U^i(s, m_i, e_j)$ represents the reward obtained by player $i (= A \text{ or } D)$ if in state s , D choose to deploy the monitoring m_i and A choose to use the exploit e_j ,
- $\gamma^i \mapsto [0, 1)$ is the discount factor for player $i (= A \text{ or } D)$.

In light of recent studies on characterizing attackers based on personality traits [2], one might argue that a defender’s viewing of rewards is different than that of the attacker’s. Given that we did not find a formal user study to show that this is the case, we will consider $\gamma^A = \gamma^D = \gamma$ going forward. As the solvers for our formulated game can work in cases even when $\gamma^A \neq \gamma^D$, this assumption just helps us simplify the notations going forward.

The concept of an optimal policy in this Markov game is well-defined for zero-sum games [17] where $U^D(s, m_i, e_j) = -U^A(s, m_i, e_j) \quad \forall s \in S, m_i \in M$, and $e_j \in E$. In these cases, a small modification to the Value Iteration algorithm can be used to compute the min-max strategy for both players. To see this, note that the Q-value update for this Markov Game (for player x) becomes as follows,

$$Q^x(s, m_i, e_j) = R^x(s, m_i, e_j) + \gamma \sum_{s'} \tau(s, m_i, e_j, s') \cdot V^D(s') \quad (1)$$

where $V^D(s')$ denotes the value function (or reward-to-go) with respect to D if in state s' . We will use the notation $M(s)$ (and $E(s)$) to denote the set of defender actions (and attacker actions) possible in state s . Given this, the mixed policy $\pi(s)$ for state s over the defender’s applicable actions ($\in M(s)$) can be computed using the value-update,

$$V^x(s) = \max_{\pi(s)} \min_{e_j} \sum_{m_i} Q^x(s, m_i, e_j) \cdot \pi_{m_i} \quad (2)$$

where π_{m_i} denotes the probability of choosing the monitoring strategy m_i .

When the Markov Game has a general-sum reward structure and one player can infer the other player’s strategy before making their move, the min-max strategy becomes sub-optimal and one must consider other notions of equilibrium [32, 8]. We give an overview of prior work on these lines later in the paper.

2.5 Quantifying the Impact of Vulnerabilities

The use of the Common Vulnerability Scoring System (CVSS) for rating the impact of attacks is well studied in security [9, 28]. For (almost all) CVEs listed

in the NVD database, we have a six-dimensional CVSS v2 vector, which can be decomposed into multiple measures like Access Complexity (AC) that model the difficulty for an attacker to exploit a particular vulnerability, and the impact on Confidentiality, Integrity, and Availability (CIA score) gained by exploiting the vulnerabilities in the cloud network above. The values of AC are categorical {EASY, MEDIUM, HIGH} (that have a corresponding numerical value associated with them), while CIA values are in the range [0, 10]. For the set of vulnerabilities present in our system, the metrics are shown in Table 1.

3 Markov Game Modeling

In this section, we model the problem of placing detection systems on the cloud as a general-sum Markov Game. An attacker’s strategy in this game is to select an attack policy based on what privilege it currently has while the defender’s strategy is to place detection mechanisms that minimize performance impact and ensure the attacker does not achieve its objectives. Thus, this work is in line with the notion of Moving Target Defense (MTD) methods for shifting the detection surface of a software system, similar to [31, 26] but we account for multi-stage attacks by using a Markov game formulation.

Beyond the obvious Markovian assumption, we further assume that (1) there is a list of attacks known to both the attacker and the defender (which cannot be immediately fixed either due to lack of manpower or restrictions from third-party customers who host their code on the cloud system [10, 26]), and (2) the attacker can be in the system and will remain undetected until it attempts to exploit an existing vulnerability, i.e. stealthy attacker [31] and (3) the state of the game is visible both to the attacker and the defender. The last assumption is often a strong one but leads to reasonably good strategies for moderately sized cloud networks. In contrary, works that consider partial observability in multi-stage settings [20, 21] find it difficult to scale beyond toy problems. In this paper, we also relax two inherent assumption of Stackelberg Security Games [23, 30, 26]— (1) detection, when a detection mechanism is in place, is perfect and (2) an attack always succeeds if it was not detected. We will later see how these uncertainties can be captured by the transition function of our Markov Game. We shall now discuss the various components of our Markov Game.

States For the Markov Game (MG) formulation, we consider the nodes N_D and N_G in the Attack Graph (AG) and map them to states in our game (see the blue diamond shaped nodes mapped to the blue circular nodes in Fig. 2). These nodes in the AG represent the state of an attacker in the cloud system, highlighting the servers they have access to and the privilege they have on these servers. Often, the location of an attacker on multiple physical servers in the cloud network can map to a single state of the AG (and therefore, a single state in our MG). The MG states that map to a goal or a root node N_G of an AG are terminal self-absorbing states, i.e. if the game reaches this state, it stays there till the end of time. Since N_G nodes represent attack goals, they correspond to states with high positive reward for the attacker in our Markov Game, representative of the fact that the attacker was successful in evading the various defense mechanisms

in place to detect attack actions. Although a zero-sum game, the defender is penalized (with a high negative reward different magnitude) for failing to detect even a single step of the attack. The other non-terminal states of the MG have a one-to-one mapping to nodes in the set N_D of the AG. Among these non-terminal states, there exist a set of states S_i the represent the external facing entry-points to the cloud network. These can be considered as the initial state of any attack.

For the cloud scenario shown in Fig 1, we have four states of which s_0 corresponding to the goal node $A-\{FTP:root\}$ is a terminal state of the MG and $S_i = \{s_1\}$ is a singleton set corresponding to the default state where an attack starts. The other two states s_2 and s_3 correspond to nodes where the attacker has different access privilege on the different three servers (LDAP, WS and FTP) server. Note that in this case, we use a ternary range to denote an adversary's privilege—no-access, user or root-user. With three servers, there can be a maximum of nine states in the AG, and hence our MG. Given the set of known attacks for our scenario, we only need to consider four of them. The maximum number of states in the MG is bounded by ($\#$ servers \times $\#$ access-levels).²

Players and Pure Strategies As mentioned above, the players for our game are the Admin (or the defender) D and the attacker A . The pure strategy set for A in state s consists of exploit actions it can perform with the privilege it has in state s . For example, in state s_2 , the attacker has the access vector $A-\{LDAP:root, FTP:user, WS:user\}$. With this as the precondition, the attack actions can be represented by the rule nodes N_R (shown in oval) in the AG. Notice that the vulnerability nodes N_C (shown as rectangles) can either be considered a part of the state s_2 or a part of the attack action. Although we don't make the distinction in this paper, we believe the latter is a better model. To see why, consider the case of multiple attackers— we might want to distinguish that an attacker a_1 , who can execute database injection exploits from another attacker a_2 , who is more capable of executing other attacks. At that point, considering the attacks as a part of the state s muddles this distinction.

The pure strategy set for D in a state s consisting of monitoring actions where each such action corresponds to placing an Intrusion Detection System (IDS) for detecting attacks that can be executed by A in state s . These actions are made possible in real-world scenarios by using two sorts of IDS systems— (1) host based IDS like *auditd* that can notify D about abnormality in the use of CPU resources or access to files on the server and (2) network based IDS like *snort* that can observe traffic on the wire and report the use of unexpected bit patterns in the header or body of a packet. Although, in the simple example, defender strategies monitor a single attack, it might be possible to take two or more detection systems and consider them together as a pure strategy for the defender. We will see this in the case of the real-world cloud scenario discussed in the experimental section where two NIDS that look for similar sorts of pattern in the network traffic can detect multiple attacks at once. In the context of Stackelberg Security Games, such groups of actions are often called schedules

² Partial observability over the state space can increase the number of states to be a power-set of this number, i.e. $2^{(\# \text{ servers} \times \# \text{ access-levels})}$

		D (Defender)		
		no-mon	mon-Web	mon-FTP
A (Attacker)	no-act	0, 0	0, -2	0, -3
	exp-Web	7, -7	-8, 6	7, -10
	exp-FTP	10, -10	10, -12	-8, 5

Table 2. Utilities (U^A, U^D) for state s_2

		no-mon	mon-LDAP			no-mon	mon-FTP
		no-act	exp-LDAP			no-act	exp-FTP
		0, 0	5, -5			0, 0	10, -10
		0, -2	-5, 3			0, -2	-8, 6

Table 3. Utilities (U^A, U^D) for states s_1 (left) and s_3 (right).

[14, 23] and the pure strategy is defined over these schedules. We note that our modeling supports such representations.³

To allow for a realistic setting, we add two other actions—**no-act** and **no-mon**—that represent a no-op for each player. This allows an attacker to not attack in a particular state if it feels that it has a high risk of getting caught. Similarly, this allows a defender to, probabilistically, not monitor for a particular attack at times, thereby saving them valuable resources.

Transitions The transitions in our MG represent that given a particular state and a pair of actions drawn from the joint action space $E \times M$, the probability with which a game reaches a state s' , i.e. $\tau(s, m, e, s')$. There exists a few obvious constraints in the context of our MG— (1) $\tau(s, m, \text{no-act}, s) = 1$, i.e. if an attacker does not execute an attack, the game remains in the same state, (2) when $e \neq \text{no-act}$, $\tau(s, m_e, e, s') = p/|S_i| \forall s' \in S_i$ where p is the probability that e is detected by m_e , the monitoring service deployed to detect e , i.e. when successfully detected, the attacker starts from either of the initial states with equal probability, and (3) $\tau(s, \text{no-mon}, e, s') \neq 0$ if $s \notin S_i$ and $s' \in S_i$, i.e. an attacker cannot be detected if the defender does not perform a monitoring action.

We highlight a few transitions for our Markov Game in Fig 2. In state s_1 , the defender does not monitor the exp-LDAP attack action and with 0.6 probability the game moves to the state s_2 (and with 0.4 it remains in s_1). These probability are calculated using the Access Complexity vector and the function defined in [5] for obtaining the probability of success given an attack. This is also done when the defender deploys a monitoring mechanism m_e but the attacker executes another attack e' where $e \neq e'$ and $e' \neq \text{no-act}$ (see the transition for an example joint action from s_2 in Fig 2). Lastly, the transition from s_3 shows a case relating to (3) mentioned in the previous paragraph. The value of $p = 0.9$ because in this case, monitoring access to files like `/etc/passwd` and `/etc/shadow` will detect a remote code execution that tries to gain root access if the attacker creates a new user or tries to escalate the privilege of a non-root user to root but may not be able to monitor a case where it simply obtains access to a root user account.

³ In these cases, the Subset of Sets are Sets (SSAS) property defined in [15] may not hold and thus, the Strong Stackelberg Equilibrium will not always be a Nash Equilibrium for the formulated Markov Game (see later (*Lemma 1*) for details).

Rewards The reward metrics for each state (except the terminal state s_0), for our small-scale example cloud scenario is shown in Table 3 and 2. Most research works [25, 21, 31] have selected random values for attacker/defender utility and cost without proper justification. The reward values in our framework are obtained using multiple metrics– (1) the impact score (IS) (also called the CIA impact score) of a particular attack, (2) the cost of performance degradation (provided by security and engineering experts in the cloud domain– often done by running MiniNET simulation) based on the placement of a particular IDS, shown experimentally by Sengupta *et. al.* [27], and (3) the hops taken by the attacker to reach a state, which is a measure of how advanced an Advanced Persistent Threat (APT) is. Note that the last factor is ideally a non-Markovian function in the overall reward estimation of our Markov Game because it depends on the path taken by the attacker to reach a particular state. To bypass this issue, we consider all possible paths an attacker can take to reach the particular state and average the path value, which gives us an average of how advanced is the APT. Given that obtaining knowledge about the actual path taken by a stealthy adversary, who has been residing in the network for a long time, is difficult (if not impossible) to obtain, as an average estimate is a pretty good heuristic for estimating the importance of an APT.

Now, let us explain the reward value for the action pair (**mon-Web**, **exp-Web**) shown in Table 2. Note that the impact score for this vulnerability (CVE-2017-5059 shown in Table 1) is 7. We utilized performance monitoring using Nagios [7] to measure the end-to-end network bandwidth, number of concurrent requests to live web services and delay in servicing network requests. We observed that there is an increase in network delay, and decrease in network bandwidth (18 Gbps to 6 Gbps as number of IDS increase from 1 to 15) [27], and number of concurrent requests serviced per second. Base on expert knowledge we estimate the reward (or rather impact on performance) of placing the IDS that monitors this vulnerability is -2 and finally this vulnerability can only be executed if the attacker has exploited at least one vulnerability before coming to this state (and thus, APT score is 1).

Thus, the defender’s reward’s for placing the correct IDs that can detect the corresponding attacker action is 7 minus 2 (cost incurred due to reduced performance) plus 1 (for detecting an APT that had already penetrated 1-hop into the network), totaling 6. On the other hand, the attacker’s reward for this action pair is -7 spending effort is executing a vulnerability of impact 7 plus -1 for losing a vantage point in the cloud network, totaling a reward of -8 . The other reward values are defined in a similar manner. Also, notice that the defender’s cost of placing IDS is not of any concern for the attacker and thus, when an attacker chooses **no-act**, their reward is 0, while the defender might still incur a negative reward for **no-mon** because it degraded the performance of the sub-net represented by the particular state.

3.1 Optimal Placement Strategy

As mentioned earlier, finding the optimal solution to a two-player general-sum Markov Games is more involved than finding the min-max strategy. Moreover, in our threat model, we assume that the attacker A will, with reconnaissance

Algorithm 1 Dynamic Programming for finding SSE in Markov Games

```

1: procedure GIVEN  $(S, M, E, \tau, U^D, U^A, \gamma^D = \gamma^A = \gamma)$ ,
2: OUTPUT  $(V^i(s), \pi^i(s) \forall i \in \{A, D\})$ 
3:    $V(s) = 0 \forall s$ 
4:   loop:  $i == k$  break;
5:   // Update Q-values
6:   Update  $Q^D(s, m, e)$  and  $Q^A(s, m, e) \forall s \in S, m \in M(s), e \in E(s)$ 
7:   using  $U^D, U^A$  and  $V(s)$ .
8:   // Do value and policy computation
9:   Calculate  $V^i(s)$  and  $\pi^i(s)$  for  $i \in \{A, D\}$  using the values  $Q^i(s, m, e)$  in Eq. 3
10:   $i \leftarrow i + 1$ 
11:  goto loop.
12: end procedure

```

efforts, get to know the strategy of the defender D , thus imparting it the flavor of leader-follower games.

We highlight a dynamic programming approach shown in Algorithm 1. This algorithm is similar to the min-max equilibrium computation method with an important change in line 9. In this step, instead of using equation 2 to calculate the optimal value and player strategies, we compute the Strong Stackelberg Equilibrium (SSE) in each state. In essence, for each iteration, when computing the value function for a particular state s , we can consider the Q-values for that state and the joint actions as a normal-form game matrix and then find the optimal policy that maximized the value of this state. Since our model, at present, does not consider multiple adversary types, the equilibrium calculation for each state can be done in polynomial time [14]. This type of a solution resembles the idea of finding Stackelberg Equilibrium in discounted stochastic game which has been discussed in [32]. They authors describe an ILP approach over all the states of the Markov Game, which becomes troublesome in our case as the number of states increase. Furthermore, the iterative approach provides an anytime solution and can be stopped at a premature stage (by setting lower values of k in Algo. 1) to yield a strategy for placement. For completeness, we now state the optimization problem in [22] that we use to update the value function for each state in each iteration of Algorithm 1.

$$\begin{aligned}
& \max_{\pi^D, \pi^A} \sum_{m \in M} \sum_{e \in E} Q^D(s, m, e) \pi_m^D \pi_e^A \tag{3} \\
& s.t. \quad \sum_{m \in M} \pi_m^D = 1, \forall \pi_m^D \pi_m^D \in [0, 1] \quad \text{Defender's selects a valid mixed strategy.} \\
& \quad \sum_{e \in E} \pi_e^A = 1, \forall \pi_e^A \pi_e^A \in \{0, 1\} \quad \text{Attacker's selects a valid pure strategy.} \\
& 0 \leq v - \sum_{m \in M} Q^A(s, m, e) \pi_m^D \leq (1 - \pi_e^A)L \quad \forall \pi_e^A \quad \text{Attacker's pure strategy maximizes their reward given defender's mixed strategy.}
\end{aligned}$$

where L is a large positive number. Given that we consider a Markov Game formulation, even with reconnaissance on an attacker's side, a threat model that

assumes that an attacker always knows the defender’s strategy in each state may be too strong. Thus, one might question the optimality of the strategy that we come up with using Algorithm 1. In [16], researchers have shown that SSEs are a subset of Nash Equilibrium for a particular class of problems. Specifically, they show that if the security resource allocation problem (which in our case, is allocating IDS for covering a vulnerability) has a particular property, termed as SSAS, then the defender’s SSE is also a NE for the game. Given this, we can state the following.⁴

Lemma 1. *If the Subset of Set Are Sets (SSAS) property holds in every state s of a Markov Game, then the SSE is also a NE of the Markov Game.*

Proof. We will prove the lemma by contradiction. Hence, let us assume that SSAS property holds for a Markov Game (MG) but the SSE is not the NE for this MG. First, consider $\gamma = 0$ for this MG. Thus, the SSE and NE strategy for each state can be calculated only based on the utilities of only this state. Now, if $\text{SSE} \not\subseteq \text{NE}$ for this Markov Game, then there is some state in which the SSE strategy is not the NE strategy. But if that is the case, then we would have violated the SSAS theorem in [16] for state, which cannot be true. For the case $\gamma > 0$, the proof still holds because note that the SSAS property is not related to the Q-values using which the strategy is computed in the Markov Game. \square

This holds trivially for our small example since the defender’s pure strategy for each state is to deploy a single IDS (and thus all subset of schedules are a possible schedule). We also seen in the upcoming experimental section 4 that the SSE and the NE for our Markov Game are the same.

4 Experimental Results

In this section, we first compare the effectiveness of the optimal strategies for our general-sum leader-follower Markov-game against the Uniform Random Strategy, which is a popular method used as a baseline for Moving Target Defense strategies in cybersecurity. We then discuss the set-up of a real-world, small scale cloud scenario and the gains we obtain using our formulated Markov Game.

4.1 Evaluation of Strategies

We first discuss two baseline strategies and then explain why we choose it for comparison with the optimal strategies obtained by solving the general-sum leader-follower markov-game formulated in the previous section.

- *Uniform Random Strategy (URS)* In this, the defender selects a pure-strategy based on a uniform random distribution. For example, in the state s_1 shown in Table 2, the defender can choose to monitor the FTP or the Web server or neither of them, all with equal probability of 33.33%. Thus, in any round, the defender rolls a three-sided fair die and picks what they should do.

Researchers have claimed that selecting between what to choose when shifting between MTD configurations should be done using a uniform random strategy

⁴ In the case of multiple attackers, $\text{SSE} \not\subseteq \text{NE}$. Although such scenarios exist in cybersecurity settings, we consider a single attacker in this modeling and plan to consider the multiple attacker setting in the future.




Fig. 3. Defender's value for each of the four state— s_1 (top-left), s_2 (top-right), s_3 (bottom-left), and s_0 , which in the all absorbing terminal state (bottom-right).

[33]. Although there have been other methods based on Stackelberg Security Games (SSGs) which have shown that such a strategy may be suboptimal [26], it provides a baseline strategy that can be used in the context of our Markov Game. Adapting the latter strategies proposed in previous work need us to compile our multi-stage markov into a single step normal form game. First, there is no trivial way of doing this conversion as the rewards in [26] talk about average impact on the network that are difficult to encode meaningfully in our Markov Game. Furthermore, attack strategies would have to represent full attack paths as opposed to single attack actions. Second, our work can be seen as a generalization of applying the notion of stackelberg equilibria, similar to [32], for Markov Games in the context of IDS placement and thus, a counterpart solution to the normal form games case described in [26] to Markov Games. Hence, we do not consider [26] as a candidate for comparison.

- *Min-max Strategy* Although our game is a general sum setting, we ask how sub-optimal are the min-max strategies for zero-sum Markov games when we ignore the impact on performance. In essence, the attacker still has the same utility in the individual states shows in Tables 3 and 2, but the defender's reward value, when calculating the strategy, is just the opposite of the attacker's reward, making this a zero-sum game by ignoring the defender's consideration for performance impacts. Once we obtain the min-max strategy, we calculate the value in each state of our general-sum Markov Game when the defender's follows the min-max policy.

Comparison of Strategies In Figure 3, we plot the values of the four states ($V(s)$) of our game for the baseline strategies (URS and Min-max), the Strong Stackelberg Eq. (SSE) and also, to provide an empirical result for Lemma 1, the Nash Eq (NE). On the x-axis, we vary the discount factor and on the y-axis plot the value function of the state for the defender. In the terminal state s_0 , the attacker has exploited all possible vulnerabilities and gains a high negative

reward. Thus, for all the states, since there is non-zero probability of reaching s_0 , the defender's value function is negative. Note that as one weighs the future rewards more, i.e. γ approaches 1, the negative value of the states increase in magnitude because the negative reward in s_0 is given higher weight.

As stated above, the SSE for our example game is the same as the NE for our Markov Game and thus, the curves for both the strategies are overlapping. On the other hand, URS is much worse off than our strategy for all the states with more than a single action whereas, Min-max, although better than URS, is sub-optimal with respect to ours for all states except s_3 and s_0 . Firstly, note that s_0 being a terminal state, has only one action for the defender. With one action, all the strategies are equivalent and we have an overlapping graph (bottom-right in Fig 3). Secondly, in state s_3 , that is just a single action away from the terminal state with high negative reward, the defender always picks the action to monitor attack traffic regardless of the performance impact (whose magnitude is less than the impact of an attack). Thus, even though the Min-max strategy is ignorant of the latter costs, it picks the same strategy as SSE. Hence, they have overlapping graphs (bottom-left in Figure 3).

Before discussing the detail about the difference between SSE and Min-max in the other two states, we take a look at the optimal mixed strategy for our game. This will help us in explaining the sub-optimality of the Min-max strategy. The optimal SSE strategy for the defender (which is also a NE strategy) for the four different states is as follows for the discount factor $\gamma = 0.8$:

$$\begin{aligned}\pi(s_0) &: \{\text{terminate}: 1.0\} \\ \pi(s_1) &: \{\text{no-mon}: 0.097, \text{mon-LDAP}: 0.903\} \\ \pi(s_2) &: \{\text{no-mon}: 0.0, \text{mon-Web}: 0.539, \text{mon-FTP}: 0.461\} \\ \pi(s_3) &: \{\text{pi_no-mon}: 0.0, \text{mon-FTP}: 1.0\}\end{aligned}$$

Note that in our example, barring the terminal state s_2 , most other states really have only one proper detection action or pure strategy because **no-mon** asks the defender to not monitor for any attacks at all and thus, we expected that these actions would have probabilities almost equal to zero (unless it has a considerable impact on performance). In the case of state s_1 , the 0.097 probability of picking that action shows that in states far away from the terminal, the defender can choose to be a little more relaxed in terms of security and pay attention to performance. On the contrary, in state s_3 an exploit action will move the game to the terminal state that has high-negative reward for the defender and thus, the defender is asked to place all attention to security. In general, this implies that an Admin D should pay more attention to security against APTs deep within the network in states closer to a critical resource and can choose to reason about performance near the entry points of the cloud system. Thus, in these states, the optimal mix-max strategy, oblivious to the performance costs, invest more monitoring resources and thus, become sub-optimal with respect to SSE.

4.2 Case Study on the Cloud

In this section, we talk about a case study on a real-world sub-network setup in a cloud system, highlighting briefly the system setup, the Markov Game formulation, the comparison between URS and SSE for a hand-picked state and how all these strategies can be enforced with the help of Software Defined Networking.

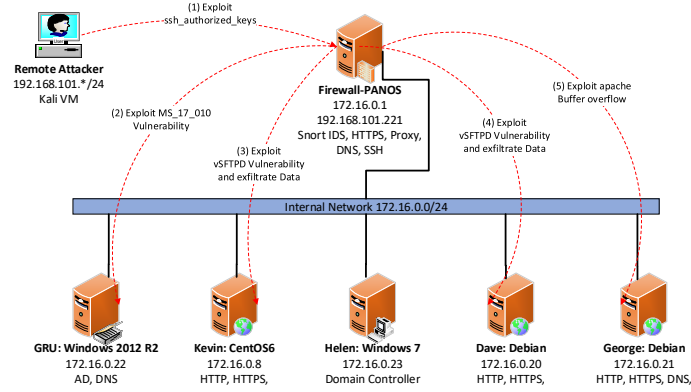


Fig. 4. A real world cloud system.

Implementation details. We utilized the Virtual Machine (VM) images from the Western Region Cybersecurity Defense Competition (WRCCDC) [6]. The competition helps university students (the Blue Team) gain first-hand experience in dealing with real-world attacks. In this competition, they are asked to defend a corporate infrastructure against experienced white-hat hackers (the Red-Team). In the scenario show in Figure 4, a Red Team attacker mounts a multi-step attack by following a slow and low approach, that tries to evade the IDS placed.

The goal of the attacker can be either the disruption of network services or ex-filtration of data out of the private networks. Both these attacks, if successful, can lead to the loss of mission-critical information (FTP server files are valuable to the company) and business (service downtime). We model each of these goal states in the attack graph as the state that leads to an all-absorbing terminal state, thus ending the game.

The set of attacks were discovered using low-intensity network scanning tools like OpenVAS that run over an extended period of time and generate a report of the vulnerabilities present in the system. Due to space considerations, we summarize the vulnerabilities present in our cloud network in Table 5. Corresponding to the attacks, we consider the deployment of IDS mechanisms like Snort IDS, Web Proxy, etc. for threat detection at different levels of the protocol stack. We used WRCCDC’s VM images to create a similar environment in our organization’s cloud service. To connect these VMs, we created a flat structure network with *Palo Alto Network OS (Next-Generation Firewall)* hosted at the gateway of the network (172.16.0.0/24) and had eight host machines in total [3].

The network was connected using SDN switch with OpenFlow v1.3 for (1) vulnerability scanning to gather knowledge about known attacks in the cloud, (2) computing the Markov Game strategy and (3) enforcing a particular deployment strategy and switching to a new one after a fixed time period T . For the first step we used scanning tools like OpenVAS, as described earlier. For the second step, we use our strategy computation algorithm that also solves the optimization problem using Gurobi solver. For the last step, we had an enable and disable (kill) script for starting and stopping the different IDS mechanisms. SDN, being

Host	High	Medium	Low
172.16.0.22	4	14	1
172.16.0.23	2	3	1
172.16.0.8	3	8	3
172.16.0.16	0	13	6
172.16.0.20	0	2	1
172.16.0.11	0	1	2
172.16.0.1	0	0	1
172.16.0.21	0	0	1
Total	8	9	41
			16

Fig. 5. The number of vulnerabilities in our cloud system found via persistent exploration with OpenVAS vulnerability scanner.




Fig. 6. Defender's value for the state s_1 as discount factor increases from 0.5 to 1.

a centralized mechanism, upon roll of a die, sent the start-stop signal to the various detection processes, which was equivalent to implementing a strategy.

Results. In the formulated Markov Game, we had eight states corresponding to each VM in our cloud system. In figure 6, we highlight the defender's value for state s_1 . This state s_1 (attacker has root privilege on Firewall VM in this state) was a lucrative vantage point for an attacker in the network because all other states were reachable via an array of vulnerabilities available on s_1 . The defender actions for this state were considered as a set that (1) detected a particular set of vulnerabilities and (2) were easy to deploy together. For example, deployment of two Network-based IDS using Snort is easier to configure for the Admin than the deployment of a host-based and a network-based IDS at the same time.

In contrast to the example scenario, in this setting, we (1) did not have a terminal state that had high negative utility and (2) the found that magnitude of positive defender utility (on successfully catching an attack) and negative defender utility (calculated using Mini-NET simulation similar to [26]) were comparable. Thus, the defender's value for state s_1 was always non-negative. As the highest gains in the value function is observed for s_1 , we plot graphs similar to the ones in the previous section for it. A key reason for URS being sub-optimal was because it neither paid more attention to attacks that had higher impact nor cared about the performance, both of which were essential given s_1 is a critical vantage point in the cloud network.

5 Related Work

Sheyner *et al* [11] present a formal analysis of attacks on a network with cost-benefit analysis and talk about potential security measures to defend against the network attacks. In [4], authors provide a polynomial time method for attack graph construction and network reconfiguration using a parallel computing approach, making it possible to gather information about attacks in large-scale systems like cloud networks. The research works, however, model the security situation from the attacker's point of view and try to characterize the qualitative impact of an attack. We consider attack-defense as a two-player game, which

provides a more realistic measure of the effectiveness of security countermeasure, and, can provide guidelines for countermeasure deployment in a cloud network.

Authors in [12] introduced the idea of moving secret proxies to new network locations using a greedy algorithm, which they show can thwart brute force and DDoS attacks. In [34], Zhuang *et al* shows that the MTD system designed with intelligent adaptations improve the effectiveness further. In [28] authors show that intelligent strategies based on common intuitions can be detrimental to security and highlight how game theoretic reasoning can alleviate the problem. On those lines, Wei *et al* [18] and Sengupta *et al* [26] use a game theoretic approach to model the attacker-defender interaction as a two-player game where they calculate the optimal response for the players using the Nash and the Stackelberg Equilibrium concepts respectively. The flavor of the approaches are similar to those of Stackelberg Security Games (SSGs) for numerous physical security applications highlighted in [22, 30]. Although they talk about the use of the Markov Decision Process (MDP) approaches for MTD, they leave it as future work. On these lines, authors in [32], show that the SSE of markov games can be arbitrarily sub-optimal for stochastic discounted path games and provide an LP approximation for the former. In this work, we believe that the Markovian assumption is sufficient to capture the strategy of an attacker and propose a dynamic programming based anytime solution method to find the SSE.

In the context of cloud systems, [24] discusses a risk-aware MTD strategy where they model the attack surface as a non-decreasing probability density function and then estimate the risk of migrating a VM to a replacement node using probabilistic inference. In [13], authors highlight obfuscation as a possible MTD strategy in order to deal with attacks like OS fingerprinting and network reconnaissance in the SDN environment. Furthermore, they highlight that the trade-off between such random mutations, which may disrupt any active services, require analysis of cost-benefits.

In this paper, we identify an adaptive MTD strategy against multi-hop monotonic attacks for cloud networks which optimizes for performance while providing gains in security. The ability to decompose a large cloud network into sub-nets provides gains in computing strategies, a fair distribution of IDS resources and prioritizing detection of attacks that may have long-term impacts.

6 Conclusion and Future Work

A cloud network is composed of heterogeneous network devices and applications interacting with each other. The interaction of these entities over a complex and hierarchical network structure poses a substantial risk to the overall security of the cloud system. At the same time, it makes the problem of monitoring ongoing attacks by adversaries located both outside and inside the network a challenging problem. In this paper, we model the concept of Moving Target Defense (MTD) for shifting the detection surface in the cloud system as a Markov Game. This helps us reason about (1) the security impact of multi-stage attacks that are characteristics of Advanced Persistent Threats (APT) while (2) ensuring that we do not place all possible security mechanisms in the cloud system, thereby hogging all the valuable cloud resources. The various parameters of our Markov Game are obtained using an array of softwares, prior research and publicly available

metrics ubiquitous in the security community. We propose a dynamic programming based anytime algorithm to find the Strong Stackelberg Equilibrium of our Markov Game and highlight its superiority to baseline strategies for an emulated and a small real-world cloud network setup.

References

1. National vulnerability database. <https://nvd.nist.gov>, accessed: 2018-09-25
2. Basak, A., Černý, J., Gutierrez, M., Curtis, S., Kamhoua, C., Jones, D., Bošanský, B., Kiekintveld, C.: An initial study of targeted personality models in the flipit game. In: International Conference on Decision and Game Theory for Security. pp. 623–636. Springer (2018)
3. Chowdhary, A., Dixit, V.H., Tiwari, N., Kyung, S., Huang, D., Ahn, G.J.: Science dmz: Sdn based secured cloud testbed. In: IEEE Conference on Network Function Virtualization and Software Defined Networks (2017)
4. Chowdhary, A., Pisharody, S., Huang, D.: Sdn based scalable mtd solution in cloud network. In: ACM Workshop on Moving Target Defense (2016)
5. Chung, C.J., Khatkar, P., Xing, T., Lee, J., Huang, D.: Nice: Network intrusion detection and countermeasure selection in virtual network systems. *IEEE transactions on dependable and secure computing* **10**(4), 198–211 (2013)
6. Competition, W.R.C.D.: WRCCDC. <https://archive.wrccdc.org/images/2018/> (2018)
7. Enterprises, N.: Nagios (2015)
8. Guerrero, D., Carsteanu, A.A., Huerta, R., Clempner, J.B.: An iterative method for solving stackelberg security games: A markov games approach. In: Electrical Engineering, Computing Science and Automatic Control (CCE), 2017 14th International Conference on. pp. 1–6. IEEE (2017)
9. Houmb, S.H., Franqueira, V.N., Engum, E.A.: Quantifying security risk level from cvss estimates of frequency and impact. *JSS* **83**(9), 1622–1634 (2010)
10. Jajodia, S., Park, N., Serra, E., Subrahmanian, V.: Share: A stackelberg honey-based adversarial reasoning engine. *ACM Transactions on Internet Technology (TOIT)* (2018)
11. Jha, S., Sheyner, O., Wing, J.: Two formal analyses of attack graphs. In: Computer Security Foundations Workshop, 2002. Proceedings. 15th IEEE. pp. 49–63. IEEE (2002)
12. Jia, Q., Sun, K., Stavrou, A.: Motag: Moving target defense against internet denial of service attacks. In: 2013 22nd International Conference on Computer Communication and Networks. pp. 1–9. IEEE (2013)
13. Kampanakis, P., Perros, H., Beyene, T.: Sdn-based solutions for moving target defense network protection. In: IEEE 15th International Symposium on a World of Wireless, Mobile and Multimedia Networks. IEEE (2014)
14. Korzhyk, D., Conitzer, V., Parr, R.: Complexity of computing optimal stackelberg strategies in security resource allocation games. In: AAAI (2010)
15. Korzhyk, D., Yin, Z., Kiekintveld, C., Conitzer, V., Tambe, M.: Stackelberg vs. Nash in Security Games: An Extended Investigation of Interchangeability, Equivalence, and Uniqueness. *J. Artif. Int. Res.* **41**(2), 297–327 (May 2011), <http://dl.acm.org/citation.cfm?id=2051237.2051246>
16. Korzhyk, D., Yin, Z., Kiekintveld, C., Conitzer, V., Tambe, M.: Stackelberg vs. nash in security games: An extended investigation of interchangeability, equivalence, and uniqueness. *Journal of Artificial Intelligence Research* **41**, 297–327 (2011)

17. Littman, M.L.: Markov games as a framework for multi-agent reinforcement learning. In: Eleventh International Conference on Machine Learning (1994)
18. Lye, K.W., Wing, J.M.: Game strategies in network security. *International Journal of Information Security* (2005)
19. McCumber, J.: Information systems security: A comprehensive model. In: Proceedings of the 14th National Computer Security Conference (1991)
20. Miehl, E., Rasouli, M., Teneketzis, D.: Optimal defense policies for partially observable spreading processes on bayesian attack graphs. In: Proceedings of the Second ACM Workshop on Moving Target Defense. pp. 67–76. ACM (2015)
21. Nguyen, T.H., Wright, M., Wellman, M.P., Singh, S.: Multistage attack graph security games: Heuristic strategies, with empirical game-theoretic analysis. *Security and Communication Networks* **2018** (2018)
22. Paruchuri, P., Pearce, J.P., Marecki, J., Tambe, M., Ordonez, F., Kraus, S.: Playing games for security: An efficient exact algorithm for solving bayesian stackelberg games. In: AAMAS, 2008. pp. 895–902 (2008)
23. Paruchuri, P., Pearce, J.P., Marecki, J., Tambe, M., Ordonez, F., Kraus, S.: Playing games for security: An efficient exact algorithm for solving bayesian stackelberg games. In: AAMAS (2008)
24. Peng, W., Li, F., Huang, C.T., Zou, X.: A moving-target defense strategy for cloud-based services with heterogeneous and dynamic attack surfaces. In: IEEE International Conference on Communications (ICC) (2014)
25. Schlenker, A., Thakoor, O., Xu, H., Fang, F., Tambe, M., Tran-Thanh, L., Vayanos, P., Vorobeychik, Y.: Deceiving cyber adversaries: A game theoretic approach. In: Proceedings of the 17th International Conference on Autonomous Agents and MultiAgent Systems. pp. 892–900. International Foundation for Autonomous Agents and Multiagent Systems (2018)
26. Sengupta, S., Chowdhary, A., Huang, D., Kambhampati, S.: Moving target defense for the placement of intrusion detection systems in the cloud. *Conference on Decision and Game Theory for Security* (2018)
27. Sengupta, S., Chowdhary, A., Huang, D., Kambhampati, S.: Moving target defense for the placement of intrusion detection systems in the cloud. In: International Conference on Decision and Game Theory for Security. pp. 326–345. Springer (2018)
28. Sengupta, S., Vadlamudi, S.G., Kambhampati, S., Doupe, A., Zhao, Z., Taguinod, M., Ahn, G.J.: A game theoretic approach to strategy generation for moving target defense in web applications. *AAMAS* (2017)
29. Shapley, L.S.: Stochastic games. *Proceedings of the national academy of sciences* **39**(10), 1095–1100 (1953)
30. Sinha, A., Nguyen, T.H., Kar, D., Brown, M., Tambe, M., Jiang, A.X.: From physical security to cybersecurity. *Journal of Cybersecurity* **1**(1), 19–35 (2015)
31. Venkatesan, S., Albanese, M., Cybenko, G., Jajodia, S.: A moving target defense approach to disrupting stealthy botnets. In: Proceedings of the 2016 ACM Workshop on Moving Target Defense. pp. 37–46. ACM (2016)
32. Vorobeychik, Y., Singh, S.: Computing stackelberg equilibria in discounted stochastic games (corrected version) (2012)
33. Zhuang, R., DeLoach, S.A., Ou, X.: Towards a theory of moving target defense. In: ACM MTD Workshop, 2014. pp. 31–40. ACM (2014)
34. Zhuang, R., Zhang, S., Bardas, A., DeLoach, S.A., Ou, X., Singhal, A.: Investigating the application of moving target defenses to network security. In: 6th International Symposium on Resilient Control Systems (ISRC). IEEE (2013)